

MPI Programming Assignment

CS 420 Fall 2022

NOTE: An automatic score of ZERO (0) will be assigned if the code fails to compile or run (fails at runtime with SEGMENTATION FAULT or something else). Points will be deducted for failing to use proper indentation (pretty printing), use of variables with meaningful names, and lack of appropriate and sufficient documentation.

1. In this homework, we will implement a simplified version of the Bully Algorithm for the **leader election problem** in distributed systems. The algorithm that you need to implement will be as follows: **[70]**

- Consider a distributed system of 'n' nodes **arranged in the form of a ring**. One of the nodes is designated as the `current_leader` at the time of MPI initialization. Your code should have the provision for changing the initial value of the `current_leader` by checking for `argv[1]`.
- Step 2: The `current_leader` can do one of the two: i) send out a HELLO message to its successor node in the ring, ii) OR decide to transmit a message to its successor in order to initiate a leader election process.
- The decision to initiate leader election has to be done in a probabilistic way. The `current_leader` generates a random number between [0, 1] and if this random number is above a certain THRESHOLD value, then initiate leader election. Leader election process involves first generating a random token (integer between 0 and MAX_TOKEN_VALUE) and then transmitting a message with tag = LEADER_ELECTION_MSG_TAG to successor. The leader election message payload consists of 2 integers – `msg[0]` = rank of the token generator, `msg[1]` = token value.
- If however, the random number is less than the THRESHOLD, then the leader would send out the message with tag=HELLO_MSG_TAG. The HELLO message consists of a single integer = HELLO_MESSAGE (declared in the header file).
- When a non-leader node 'i' receives a leader election message from its predecessor, it too first has to take a probabilistic decision as to whether it will participate in the leader election process (in exactly the same way as the `current_leader`). If yes, then it too generates a random token value (`mytoken`) and compares with the value of the token it received. if (`mytoken > msg[1]`) OR if (`mytoken == msg[1]`) AND (`myrank > msg[0]`) then the node 'i' will update the `msg[0]` and `msg[1]` with its rank and `mytoken` values respectively. If no, then it keeps the leader election message unchanged. Regardless of the outcome, node 'i' will transmit the LEADER_ELECTION_MSG to its successor. It will then issue a blocking `MPI_Recv()` call to receive the results of the leader election process by waiting for a message with tag=LEADER_ELECTION_RESULT_MSG_TAG.
- If a node receives a HELLO message, it first generates a random number and tests whether it is larger than a given TX_PROB. TX_PROB may be modified by the user at runtime by specifying `argv[2]`. If larger than TX_PROB, then it will simply send the HELLO message to its successor node. Otherwise, it will not transmit the HELLO message to its successor. This is a common way of simulating node/link failure in distributed systems design.
- Each non-leader node in the system, waits to receive either a HELLO or LEADER_ELECTION_MSG up to a TIME_OUT_INTERVAL period (in secs) using the `MPI_lprobe()` function. If no message is received within that time, then it gives up and goes for the next round → hits and `MPI_Barrier()`

statement and waits for all the other nodes in the system to also reach the barrier. This is a standard practice of implementing a round-based protocol wherein the nodes are synchronized at the beginning/end of every round. The number of rounds the algorithm will run for may be changed through user input (using `argv[2]`).

- The `current_leader` will wait to receive back the HELLO or LEADER_ELECTION_MSG using a combination of `MPI_Irecv()` and `MPI_Test()` functions. If the `current_leader` times out (due to not receiving a HELLO message), then it will cancel the `MPI_Irecv()` and then free up the receive call Request resources using `MPI_Cancel()` and `MPI_Request_free()` functions.
- NOTE: only the HELLO message may be lost in the network!
- When the `current_leader` receives back the LEADER_ELECTION_MSG, it will update its `current_leader = msg[0]` and then send out LEADER_ELECTION_RESULT_MSG to its successor which is a 2-element integer array consisting of: `msg[0]` = new leader ID, `msg[1]` = new leader's token value. It will then also issue a blocking `MPI_Recv()` call to receive back the message with the LEADER_ELECTION_RESULT_MSG_TAG.
- When a node receives a message with the LEADER_ELECTION_RESULT_MSG_TAG, it updates its `current_leader = msg[0]` of the received message. It then prints out its new leader.
- At the end of the above steps, each node issues a `MPI_Barrier()` to allow for synchronization before starting the next round of iteration.
- To give a sense of how your program is executing, you should print out appropriate messages for each node in every round whenever it receives/sends a message; decision to participate in leader election; decision to transmit or not a HELLO message; time-out occurs; result of leader election process.

I will be providing you with skeletal file (**simplebully.h**) to help you get started as well as maintain standard notations.

Your code should take as command-line argument the number of rounds for which to run the leader election algorithm. The variable `MAX_ROUNDS` in `simplebully.h` header should be used to set the number of rounds. I would suggest starting with `MAX_ROUNDS = 1` and then once you are confident of the code, increase the number of rounds to something like `MAX_ROUNDS = 10`.

BONUS POINTS: Use `getopt()` to enable command-line input parameters. We will discuss this more in class on what parameters to enable for your code. **[+5]**

2. We discussed in class the topological sort algorithm, which plays an important role in identifying cycles in a WFG. Your objective will be to implement the topological sort algorithm by reading/creating a WFG as follows:

- i. The number of nodes 'N' in the graph will be obtained from the command-line argument (first argument to your program). The node IDs will be expected to be in the range 0 to N-1.
- ii. A filename representing a file that will contain the list of edges passed as the 2nd command-line argument. I would suggest a name like **edgelist.txt** which has the following format: Two integers separated by a whitespace per row, with the first integer representing the origin node and the second integer indicating the destination node.

As discussed in class, you should first create an adjacency matrix from the edgelist.txt file and then implement your topological sort algorithm using the constructed adjacency matrix. Your program should output either “Contains cycle”, or “Does not contain cycle” for the input WFG. **[30]**

BONUS POINTS:

- i. Use getopt() to enable command-line input parameters. We will discuss this more in class on what parameters to enable for your code. **[+5]**
- ii. Create a visualization the WFG along with the output from the topological sort. **[+15]**